

Ex5 : Implement MC Prediction in Blackjack game

```
import gymnasium as gym
```

```
from collections import defaultdict
```

```
env=gym.make('Blackjack-v1',render_mode='human')
```

1. A state is represented as a tuple of three values:

1. The value of the sum of our cards
2. The face value of one of the dealer's card
3. Boolean value — True if we have a useable ace and False if we don't have a useable ace

```
print(env.reset())
```

Output : (15, 9, True)

2. Action Space :

- The action **stand** is represented by 0
- The action **hit** is represented by 1

```
print(env.action_space)
```

Output : Discrete(2)

3. Reward :

- +1.0 reward if we win the game
 - -1.0 reward if we lose the game
 - 0 reward if the game is a draw
-

STEP 1 : DEFINE A POLICY TO BE EVALUATED

```
def policy(state):
```

```
    return 0 if state[0]>19 else 1
```

```
#generate an initial state
```

```
state = env.reset()
```

```
state=state[0]
```

```
print(state)
```

```
#print the policy of this state
```

```
print(policy(state))
```

STEP 2 : GENERATE AN EPISODE USING THIS POLICY

```
num_timesteps = 100
def generate_episode(policy):
    episode = []
    state = env.reset()
    state=state[0]
    for t in range(num_timesteps):
        action = policy(state)
        next_state, reward, done, info,trans_prob = env.step(action)
        episode.append((state, action, reward))
        if done:
            break
        state=next_state
    return episode
```

print(generate_episode(policy))

```
[((10, 2, False), 1, 0), ((20, 2, False), 0, 1.0)]
```

There are 2 states in this episode. We performed action 1 (hit) in state (10,2,False) and got a reward of 0. We performed action 0 (stand) in state (20,2, False) and got a reward of 1.0.

STEP 3: COMPUTE THE VALUE FUNCTION OF THE STATES

To predict the value function, we generate several episodes using the given policy and compute the value of the state as an average return across several episodes

```
total_return = defaultdict(float)
```

```
N = defaultdict(int)
```

```
num_iterations = 50000
```

```
for i in range(num_iterations):
```

```
    episode = generate_episode(policy)
```

```
    states, actions, rewards = zip(*episode)
```

```
    for t, state in enumerate(states):
```

```
        if state not in states[0:t]:
```

```
            R = (sum(rewards[t:]))
```

```
            total_return[state] = total_return[state] + R
```

```
            N[state] = N[state] + 1
```

```
print(total_return[state])
```

```
print(N[state])
```

```
# convert both the dictionaries to a Data frame, Merge the frames based on  
# 'state' column
```

```
total_return = pd.DataFrame(total_return.items(), columns=['state',  
'total_return'])
```

```
N = pd.DataFrame(N.items(), columns=['state', 'N'])
```

```
df = pd.merge(total_return, N, on="state")
```

```
print(df.head(10))
```

	state	total_return	N
0	(7, 10, False)	-52.0	98
1	(11, 10, False)	6.0	190
2	(15, 10, False)	-219.0	394
3	(12, 10, False)	-160.0	331
4	(18, 10, False)	-269.0	406
5	(14, 10, True)	-21.0	49
6	(21, 10, True)	176.0	193
7	(16, 4, False)	-55.0	83
8	(10, 10, False)	-23.0	150
9	(20, 10, False)	275.0	585

Figure 4.19: The total return and the number of times a state has been visited

Next, we can compute the value of the state as the average return

$$V(s) = \frac{\text{total_return}(s)}{N(s)}$$

```
df['value'] = df['total_return']/df['N']
print(df.head(10))
print(df.shape)
```

	state	total_return	N	value
0	(7, 10, False)	-52.0	98	-0.530612
1	(11, 10, False)	6.0	190	0.031579
2	(15, 10, False)	-219.0	394	-0.555838
3	(12, 10, False)	-160.0	331	-0.483384
4	(18, 10, False)	-269.0	406	-0.662562
5	(14, 10, True)	-21.0	49	-0.428571
6	(21, 10, True)	176.0	193	0.911917
7	(16, 4, False)	-55.0	83	-0.662651
8	(10, 10, False)	-23.0	150	-0.153333
9	(20, 10, False)	275.0	585	0.470085

Figure 4.20: The value is calculated as the average of the return of each state

We have successfully predicted the value function of the given policy using the every-visit MC method.

Check the value of some states using the policy we set in the beginning (which was an optimal policy), which selects the action 0 (stand) when the sum value is greater than 19 and the action 1 (hit) when the sum value is lower than 19.

Question 1 : Evaluate the value of the state (21,9,False). This is a good state and hence should have a high value.

```
df[df['state']==(21,9,False)]['value'].values
```

Output :

```
array([1.0])
```

Question 2 : Evaluate the value of the state (5,8,False). This is not a good state and hence should have a low value.

```
df[df['state']==(5,8,False)]['value'].values
```

Output :

```
array([-1.0])
```

Thus, we learnt how to predict the value function of the given policy using the every-visit MC prediction method.